

# CS 660: Mathematical Foundations of Analytics

Dr. Francis Parisi

Pace University

Spring 2018

---

## *Time Series Analysis*

- ▶ Time Series and Stochastic Processes Basics
- ▶ Stationary Processes
- ▶ ARMA Models
- ▶ Nonstationary and Seasonal Time Series
- ▶ Multivariate Time Series
- ▶ EWMA, ARCH and GARCH





# Time Series Analysis

## Time Series and Stochastic Processes Basics

1. A **time series** is a sequence of data recorded in time
2. Examples of time series
  - ▶ Australian Red Wine Sales
  - ▶ Monthly Accidental Deaths
  - ▶ U.S. Interest Rates
  - ▶ Stock Prices

# Time Series Analysis

## Examples of time series



# Time Series Analysis

Time series may

- ▶ Be random noise
- ▶ Have a trend
- ▶ Have a seasonal component
- ▶ Have both a trend and a seasonal component

# Time Series Analysis

- ▶ Modeling for time series is conceptually similar to other modeling problems
- ▶ One major distinction is that usually the next value of the series is highly related to the most recent values
- ▶ This dependence decays with time

# Time Series Analysis

- ▶ When we model a time series we first try simplify it as much as possible
- ▶ If the series has properties that change over time (like the mean or the variance), we preprocess the series to make it stationary (statistical properties constant in time)
- ▶ First we note that often there are both trends and seasonality in many data sets
- ▶ Linear trends are removed by doing a first difference on the time series, where one models the differences between the values at each time step instead of the value at each time step

# Time Series Analysis

- ▶ Seasonality can be identified and removed as follows
- ▶ First one identifies the natural periodicity of the data. This can be done in a variety of ways, such as
  - (1) through expert understanding of the dynamics
  - (2) through statistical analysis using different window lengths or
  - (3) through frequency analysis looking for the fundamental frequency of frequencies
- ▶ Once we identify the seasonal component we subtract it from the data



# Time Series Analysis

- ▶ Most time series modeling handles many of these processes via the Autoregressive Integrated Moving Average (ARIMA) methodology
- ▶ The partial autocorrelation function explicitly accounts for the simple statistical linear correlation to the past values
- ▶ The ARIMA model training is to learn the not so obvious relationships
- ▶ The moving average process helps mitigate noise

# Time Series Analysis

- ▶ The ARIMA methodology is inherently linear
- ▶ Most mainline forecasting limits itself to linear modeling
- ▶ Much like the widespread use of linear regressions
- ▶ The Generalized Autoregressive Conditional Heteroscedastic (GARCH) methodology is essentially built around autocorrelation and a moving average

# Time Series Analysis

- ▶ We often deal with multiple time series problems
- ▶ One simultaneously predicts the output of multiple time series using all the time series as inputs
- ▶ This category of problems is a straightforward extension of the single time series prediction methodology
- ▶ Multiple time series are the inputs to each of the single prediction models for each separate time series
- ▶ This is known as multivariate time series analysis, also vector autoregressive or VAR

# Time Series Analysis

General form of a time series is

$$X_t = m_t + S_t + Y_t \quad (7.1)$$

where  $m_t$  is the trend,  $S_t$  is the seasonal component, and  $Y_t$  is a mean 0 **stationary** series

Our goal: Estimate  $m_t, S_t$  and eliminate them to get a mean 0, stationary series

# Time Series Analysis

## Definition

Let  $\{X_t\}$  be a time series with  $\mathbb{E} X_t^2 < \infty$ . The **mean function** of  $\{X_t\}$  is

$$\mu_X(t) = \mathbb{E}(X_t)$$

The **covariance function** of  $\{X_t\}$  is

$$\gamma_X(r, s) = \text{Cov}(X_r, X_s) = \mathbb{E}[(X_r - \mu_X(r))(X_s - \mu_X(s))]$$

for all  $r, s, t \in \mathbb{Z}$ .

# Time Series Analysis

## Definition (2)

$\{X_t\}$  is **(weakly) stationary** if

(i)  $\mu_X(t)$  is independent of  $t$   
and

(ii)  $\gamma_X(t+h, t)$  is independent of  $t$  for each  $h$

- ▶ **Strict stationarity** of a time series  $\{X_t, t = 0, \pm 1, \dots\}$  is defined by the condition that  $(X_1, \dots, X_n)$  and  $(X_{1+h}, \dots, X_{n+h})$  have the same joint distributions for all integers  $h$  and  $n > 0$
- ▶ It is easy to check that if  $\{X_t\}$  is strictly stationary and  $\mathbb{E} X_t^2 < \infty$  for all  $t$ , then  $\{X_t\}$  is also weakly stationary

# Time Series Analysis

Given condition (ii), whenever we use covariance function with reference to a stationary series, we'll mean the function  $\gamma_X$  of one variable defined by

$$\gamma_X(h) := \gamma_X(h, 0) = \gamma_X(t + h, t)$$

The function  $\gamma_X(h)$  is the autocovariance function at lag  $h$

# Time Series Analysis

## Definition

Let  $\{X_t\}$  be a stationary time series. the **autocovariance function** (ACVF) of  $\{X_t\}$  is

$$\gamma_X(h) = \text{Cov}(X_{t+h}, X_t)$$

The **autocorrelation function** (ACF) of  $\{X_t\}$  is

$$\rho_X(h) \equiv \frac{\gamma_X(h)}{\gamma_X(0)} = \text{Cor}(X_{t+h}, X_t)$$



# Time Series Analysis

## *iid* Noise

If  $\{X_t\}$  is *iid* noise and  $\mathbb{E}(X_t^2) = \sigma^2 < \infty$ , then  $\mathbb{E}(X_t) = 0$  for all  $t$ .  
By the assumed independence,

$$\gamma_X(t+h, h) = \begin{cases} \sigma^2, & \text{if } h = 0, \\ 0, & \text{if } h \neq 0 \end{cases} \quad (7.2)$$

which does not depend on  $t$

Therefore *iid* noise with finite second moment is stationary and we denote it as

$$\{X_t\} \sim \text{IID}(0, \sigma^2)$$

to indicate that the random variables  $X_t$  are independent and identically distributed random variables, each with mean 0 and variance  $\sigma^2$ .

# Time Series Analysis

## White Noise

If  $\{X_t\}$  is a sequence of uncorrelated random variables, each with zero mean and variance  $\sigma^2$ , then  $\{X_t\}$  is stationary with the same covariance function as the *iid* noise

Such a sequence is referred to as **white noise** with mean 0 and variance  $\sigma^2$

This is indicated by the notation

$$\{X_t\} \sim \text{WN}(0, \sigma^2)$$

Every  $\text{IID}(0, \sigma^2)$  sequence is  $\text{WN}(0, \sigma^2)$  but not the converse

# Time Series Analysis

## First-order Moving Average – $MA(1)$ process

Consider the series defined by the equation

$$X_t = Z_t + \theta Z_{t1}, t = 0, \pm 1, \dots$$

where  $\{Z_t\} \sim \text{WN}(0, \sigma^2)$  and  $\theta$  is a real-valued constant

We see that  $\mathbb{E} X_t = 0, \mathbb{E} X_t^2 = \sigma^2(1 + \theta^2) < \infty$ , and

$$\gamma_X(h + t, t) = \begin{cases} \sigma^2(1 + \theta^2), & \text{if } h = 0, \\ \sigma^2\theta, & \text{if } h = \pm 1, \\ 0, & \text{if } |h| > 1. \end{cases}$$

# Time Series Analysis

Definition 2 is satisfied, and  $\{X_t\}$  is stationary

The autocorrelation function of  $\{X_t\}$  is

$$\rho_X(h) = \begin{cases} 1, & \text{if } h = 0, \\ \theta/(1 + \theta^2), & \text{if } h = \pm 1, \\ 0, & \text{if } |h| > 1. \end{cases}$$

# Time Series Analysis

## First-order Autoregression – $AR(1)$ process

Let's assume now that  $\{X_t\}$  is a stationary series satisfying the equation

$$X_t = \phi X_{t-1} + Z_t, t = 0, \pm 1, \dots$$

where  $\{Z_t\} \sim \text{WN}(0, \sigma^2)$ ,  $|\phi| < 1$ , and  $Z_t$  is uncorrelated with  $X_s$  for each  $s < t$ .

By taking expectations on each side and using the fact that  $\mathbb{E} Z_t = 0$ , we see that  $\mathbb{E} X_t = 0$

# Time Series Analysis

To find the autocorrelation function of  $\{X_t\}$  we multiply each side of the equation by  $X_{t-h}$ , ( $h > 0$ ) and then take expectations, lastly observing that  $\gamma(h) = \gamma(-h)$  we get

$$\rho_X(h) = \frac{\gamma_X(h)}{\gamma_X(0)} = \phi^{|h|}, \quad h = 0, \pm 1, \dots$$

It follows from the linearity of the covariance function in each of its arguments and the fact that  $Z_t$  is uncorrelated with  $X_{t-1}$ , that

$$\gamma_X(0) = \text{Cov}(X_t, X_t) = \text{Cov}(\phi X_{t-1} + Z_t, \phi X_{t-1} + Z_t) = \phi^2 \gamma_X(0) + \sigma^2$$

and so

$$\gamma_X(0) = \frac{\sigma^2}{1 - \phi^2}$$

# Time Series Analysis

## Sample ACF

- ▶ Normally we don't start with a model but with data
- ▶ To determine the degree of dependence among the observed data we use the sample ACF
- ▶ The sample ACF provides an estimate of the ACF of  $\{X_t\}$

# Time Series Analysis

## Definition

Let  $x_1, \dots, x_n$  be observations of a time series. The **sample mean** of  $x_1, \dots, x_n$  is

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i.$$

The **sample autocovariance function** is

$$\hat{\gamma}(h) := n^{-1} \sum_{i=1}^{n-|h|} (x_{t+|h|} - \bar{x})(x_t - \bar{x}), \quad -n < h < n.$$


The **sample autocorrelation function** is

$$\hat{\rho} = \frac{\hat{\gamma}(h)}{\hat{\gamma}(0)}, \quad -n < h < n.$$



# Time Series Analysis

ACF Plot 200 Standard Normal RV




# Time Series Analysis

- ▶ For data containing a trend,  $|\hat{\rho}(h)|$  will exhibit slow decay as  $h$  increases
- ▶ For data with a substantial deterministic periodic component,  $|\hat{\rho}(h)|$  will exhibit similar behavior with the same periodicity

# Time Series Analysis

ACF of Wine Data



# Time Series Analysis

- ▶ Recall the general form of our time series (7.1) with a trend, seasonal component and noise
- ▶ To analyze the time series we first need to plot the data and inspect the series
- ▶ If the series displays a trend and or seasonal component, we transform the data before further modeling

# Time Series Analysis

A common way to de-trend a series without seasonality is by differencing

We define the lag-1 difference operator  $\nabla$  by

$$\nabla X_t = X_t - X_{t-1} = (1 - B)X_t$$

where  $B$  is the backard shift operator

$$BX_t = X_{t-1}$$

# Time Series Analysis

Powers of the operators  $B$  and  $\nabla$  are defined in the obvious way, i.e.,  $B^j(X_t) = X_{t-j}$  and  $\nabla^j(X_t) = \nabla(\nabla^{j-1}(X_t))$ ,  $j \geq 1$ , with  $\nabla^0(X_t) = X_t$

Polynomials in  $B$  and  $\nabla$  are manipulated in precisely the same way as polynomial functions of real variables

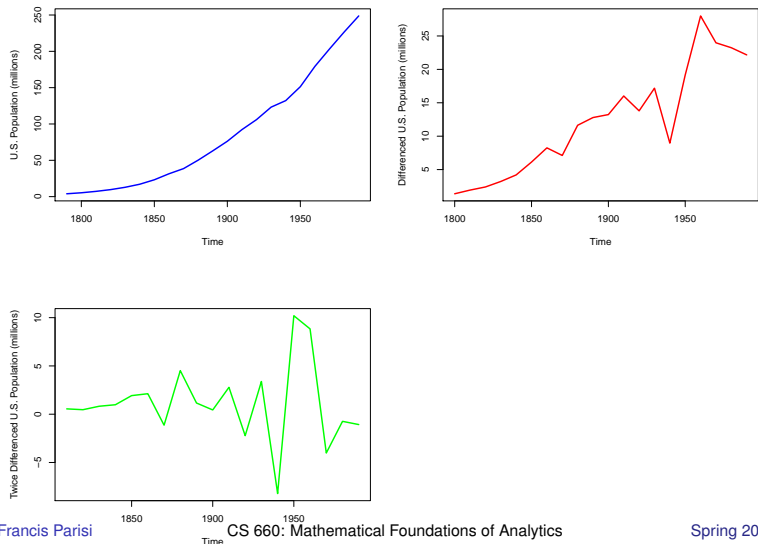
For example,

$$\begin{aligned}\nabla^2 X_t &= \nabla(\nabla(X_t)) = (1 - B)(1 - B)X_t = (1 - 2B + B^2)X_t \\ &= X_t - 2X_{t-1} + X_{t-2}\end{aligned}$$

is how we would create a twice-differenced series

# Time Series Analysis

Figure 7.1: U.S. Population data twice-differenced to remove trend



# Time Series Analysis

## Stationary Processes

Earlier we introduced and defined the **autocovariance function**  $\gamma(h)$ , and the **autocorrelation function**  $\rho(h)$

**Basic Properties of  $\gamma(\cdot)$ :**

$$\gamma(0) \geq 0,$$

$$|\gamma(h)| \leq \gamma(0) \text{ for all } h,$$

and  $\gamma(\cdot)$  is even, i.e.,

$$\gamma(h) = \gamma(-h) \text{ for all } h$$



# Time Series Analysis

We also talked about (weakly) stationary time series and now look at **strictly** stationary time series

$\{X_t\}$  is a **strictly stationary time series** if

$$(X_1, \dots, X_n)' \stackrel{d}{=} (X_{1+h}, \dots, X_{n+h})'$$

for all integers  $h$  and  $n \geq 1$ . (The symbol  $\stackrel{d}{=}$  means the two random vectors are equal in distribution, or have the same joint distribution function)

# Time Series Analysis

## Properties of a Strictly Stationary Time Series

- (a) The random variables  $X_t$  are identically distributed
- (b)  $(X_t, \dots, X_{t+h})' \stackrel{d}{=} (X_1, \dots, X_{1+h})'$  for all integers  $t$  and  $h$
- (c)  $\{X_t\}$  is weakly stationary if  $\mathbb{E} X_t^2 < \infty$  for all  $t$
- (d) Weak stationary does *not* imply strict stationary (but the other way around)
- (e) An *iid* sequence is strictly stationary

# Time Series Analysis

## ARMA Models

We've looked at  $\text{AR}(p)$  and  $\text{MA}(q)$  processes and now consider the  $\text{ARMA}(p, q)$  process

### Definition

The time series  $\{X_t\}$  is an **ARMA(1, 1) process** if it is stationary and satisfies the following for all  $t$

$$X_t - \phi X_{t-1} = Z_t + \theta Z_{t-1}, \quad (7.3)$$

where  $\{Z_t\} \sim \text{WN}(0, \sigma^2)$  and  $\phi + \theta \neq 0$ .

# Time Series Analysis

Using the backward shift operator  $B$ , we can rewrite (7.3) as

$$\phi(B)X_t = \theta(B)Z_t$$

where  $\phi(B)$  and  $\theta(B)$  are the linear filters

$$\phi(B) = 1 - \phi B \text{ and } \theta(B) = 1 + \theta B,$$

respectively.

# Time Series Analysis

Extending to  $p, q > 1$

## Definition

$\{X_t\}$  is an ARMA( $p, q$ ) process if  $\{X_t\}$  is stationary and if for every  $t$ ,

$$X_t - \phi_1 X_{t-1} - \cdots - \phi_p X_{t-p} = Z_t + \theta_1 Z_{t-1} + \cdots + \theta_q Z_{t-q}, \quad (7.4)$$

where  $\{Z_t\} \sim \text{WN}(0, \sigma^2)$  and the polynomials  $(1 - \phi_1 z - \cdots - \phi_p z^p)$  and  $(1 + \theta_1 z + \cdots + \theta_q z^q)$  have no common factors

# Time Series Analysis

The process  $\{X_t\}$  is said to be an  $\text{ARMA}(p, q)$  process with mean  $\mu$  if  $\{X_t - \mu\}$  is an  $\text{ARMA}(p, q)$  process  
It is convenient to use the more concise form

$$\phi(B)X_t = \theta(B)Z_t$$

where the  $\phi(\cdot)$  and  $\theta(\cdot)$  are the  $p^{\text{th}}$  and  $q^{\text{th}}$  degree polynomials

$$\phi(z) = 1 - \phi_1 z - \cdots - \phi_p z^p$$

and

$$\theta(z) = 1 + \theta_1 z + \cdots + \theta_q z^q$$

# Time Series Analysis

## Existence and Uniqueness

A stationary solution  $\{X_t\}$  of (7.4) exists (and is also the unique stationary solution) if and only if

$$\phi(z) = 1 - \phi_1 z - \cdots - \phi_p z^p \neq 0, \text{ for all } |z| = 1$$

# Time Series Analysis

## Causality

An ARMA( $p, q$ ) process  $\{X_t\}$  is **causal**, or a **causal function** of  $\{Z_t\}$ , if there exist constants  $\{\psi_j\}$  such that  $\sum_{j=0}^{\infty} |\psi_j| < \infty$  and

$$X_t = \sum_{j=0}^{\infty} \psi_j Z_{t-j}, \quad \text{for all } t.$$

Causality is equivalent to the condition

$$\phi(z) = 1 - \phi_1 z - \cdots - \phi_p z^p \neq 0, \quad \text{for all } |z| \leq 1$$

We can write a causal model as an infinite-order moving average model MA( $\infty$ )



# Time Series Analysis

## Invertibility

An ARMA( $p, q$ ) process  $\{X_t\}$  is **invertible** if there exist constants  $\{\pi_j\}$  such that  $\sum_{j=0}^{\infty} |\pi_j| < \infty$  and

$$Z_t = \sum_{j=0}^{\infty} \pi_j X_{t-j} \text{ for all } t.$$

Invertibility is equivalent to the condition

$$\theta(z) = 1 + \theta_1 z + \cdots + \theta_q z^q \neq 0 \text{ for all } |z| \leq 1$$

We can write a invertible model as an infinite-order autoregressive model AR( $\infty$ )

# Time Series Analysis


Earlier we introduced the autocorrelation function

The **partial** ACF measures the correlation between the residuals of  $X_t$  and  $X_{t+k}$  after the regressions on  $X_{t+1}, \dots, X_{t+k-1}$

The PACF of a causal  $AR(p)$  process is zero for lags greater than  $p$  so the PACF is helpful in identifying the order of the process


# Time Series Analysis

```
ts.plot(sunspot.year, col="blue", lwd=2)
```




# Time Series Analysis

```
acf(sunspot.year, main=" ")
```



# Time Series Analysis

```
acf(sunspot.year, type = "partial", main="")
```



# Time Series Analysis

```
sunspot.ar <- ar(sunspot.year, order.max = 2)
sunspot.pred <- predict(sunspot.ar, n.ahead = 10)


ylow = min(sunspot.pred$pred-1.96*sunspot.pred$se)
yhigh = max(sunspot.pred$pred+1.96*sunspot.pred$se)

plot(sunspot.pred$pred, type="l", col="red", lwd=2,
      ylim = c(ylow, yhigh))

points(sunspot.pred$pred+1.96*sunspot.pred$se,
       col="blue", pch=16)

points(sunspot.pred$pred-1.96*sunspot.pred$se,
       col="blue", pch=16)
```

# Time Series Analysis



# Time Series Analysis

## Deriving the Autocovariance for an AR(1)

Recall the AR(1) model has the form

$$X_t = \phi X_{t-1} + Z_t$$

where  $\{Z_t\} \sim WN(0, \sigma^2)$

We can write this as

$$X_t = \phi(\phi X_{t-2} + Z_{t-1}) + Z_t = \phi^2 X_{t-2} + \phi Z_{t-1} + Z_t$$

Recursive substitution leads to

$$X_t = Z_t + \phi Z_{t-1} + \phi^2 Z_{t-2} + \phi^3 Z_{t-3} + \cdots$$



# Time Series Analysis

Taking expectations on both sides we get  $\mathbb{E} X_t = 0$

$$\begin{aligned}\gamma(h) &= \text{Cov}(X_t, X_{t-h}) = \mathbb{E}[X_t X_{t-h}] - \mathbb{E}[X_t] \mathbb{E}[X_{t-h}] \\ &= \mathbb{E}[X_t X_{t-h}] \text{ since } \mathbb{E} X_t = 0, \forall t \\ &= \mathbb{E}[(\phi X_{t-1} + Z_t) X_{t-h}] = \phi \mathbb{E}[X_{t-1} X_{t-h}] \\ &= \phi \mathbb{E}[(\phi X_{t-2} + Z_{t-1}) X_{t-h}] = \phi^2 \mathbb{E}[X_{t-2} X_{t-h}] \\ &\vdots \\ &= \phi^h \mathbb{E}[X_{t-h} X_{t-h}] = \phi^h \text{Var}(X_{t-h}) = \phi^h \frac{\sigma^2}{1 - \phi^2}\end{aligned}$$

# Time Series Analysis

## Nonstationary and Seasonal Time Series

- ▶ Most real world time series are not stationary Oftentimes, differencing the data can help to achieve stationarity This leads to the class of autoregressive integrated moving average models, or ARIMA

# Time Series Analysis

## Definition

If  $d$  is a nonnegative integer, then  $\{X_t\}$  is an  $\text{ARIMA}(p, d, q)$  process if  $Y_t := (1 - B)^d X_t$  is a causal  $\text{ARMA}(p, q)$  process.

General form of the difference equation is

$$\phi * (B)X_t \equiv \phi(B)(1 - B)^d X_t = \theta(B)Z_t,$$

where  $Z_t \sim \text{WN}(0, \sigma^2)$

# Time Series Analysis

## Example

Suppose  $\{X_t\}$  is ARIMA(1, 1, 0)

$$(1 - \phi B)(1 - B)X_t = Z_t$$

which gives us  $(1 - \phi B - B + \phi B^2)X_t = Z_t$  or

$$X_t - \phi X_{t-1} - X_{t-1} + \phi X_{t-2} = Z_t \quad (7.5)$$

If we let  $Y_t = X_t - X_{t-1}$  then (7.5) becomes

$$Y_t = \phi Y_{t-1} + Z_t$$

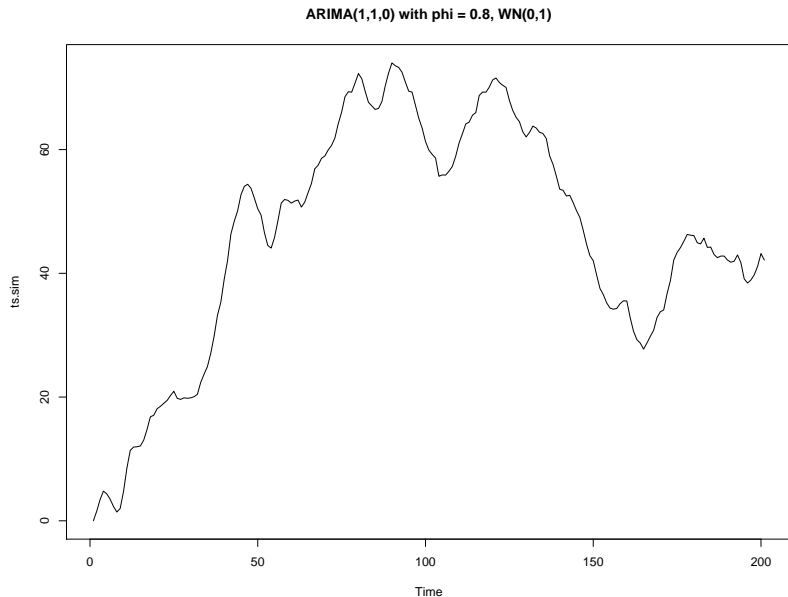
which we recognize as an AR(1) model

# Time Series Analysis

**Simulate an ARIMA(1,1,0) process, and plot the data, ACF and PACF**


```
ts.sim <- arima.sim(list(order = c(1,1,0),  
                        ar = 0.8), n = 200, sd=1)  
  
ts.plot(ts.sim,  
        main="ARIMA(1,1,0) with phi = 0.8,  
        WN(0,1) ")  
  
par(mfrow=c(1,2))  
tmp <- acf(ts.sim, main="ACF ARIMA(1,1,0) ")  
tmp <- acf(ts.sim, type="partial",  
          main="PACF ARIMA(1,1,0) ")  
par(mfrow=c(1,1))
```

# Time Series Analysis




# Time Series Analysis

ACF ARIMA(1,1,0)



PACF ARIMA(1,1,0)




# Time Series Analysis

## Fitting an ARIMA model to the data

```
#fitting an ARIMA model to data  
arima.out <- arima(ts.sim, order=c(1,1,0))  
tsdiag(arima.out)
```



# Time Series Analysis



# Time Series Analysis

## Dealing with Non-stationarity

- ▶ Visualization – plot the data
- ▶ Classical decomposition breaks out trend, seasonal, and random components
- ▶ Differencing the data to remove trend and seasonality

# Time Series Analysis

- ▶ Inspect the ACF and PACF for estimating the order(s) of the model
- ▶ Try several models and select the one with the lowest Akaike Information Criterion Corrected, or AICC

$$AICC = -2 \log \mathcal{L} + \frac{2(p + q + 1)n}{n - p - q - 2}$$

- ▶ Once we fit the model we need to check the assumptions such as white noise residuals
- ▶ Try to keep  $p + q$  reasonable; avoid large models

# Time Series Analysis

## Multivariate Time Series

- ▶ Some time series are best treated as components of a vector valued (or multivariate) time series,  $\{\mathbf{X}_t\}$
- ▶ In such cases, we have both serial dependence and interdependence between series
- ▶ The mean vector is given by

$$\boldsymbol{\mu} = \mathbb{E} \mathbf{X}$$

and the covariance matrix by

$$\Gamma(h) = \mathbb{E}[\mathbf{X}_{t+h}\mathbf{X}_t^T] - \boldsymbol{\mu}\boldsymbol{\mu}^T$$

# Time Series Analysis

- ▶ The problem of identifying the MV ARMA model can be avoided if we consider the MV AR model, or the Vector AR (VAR) model
- ▶ For a bivariate time series we have

$$\boldsymbol{\mu} = \mathbb{E} \mathbf{X} = \begin{bmatrix} \mathbb{E} X_{t1} \\ \mathbb{E} X_{t2} \end{bmatrix}$$

and the covariance matrix by

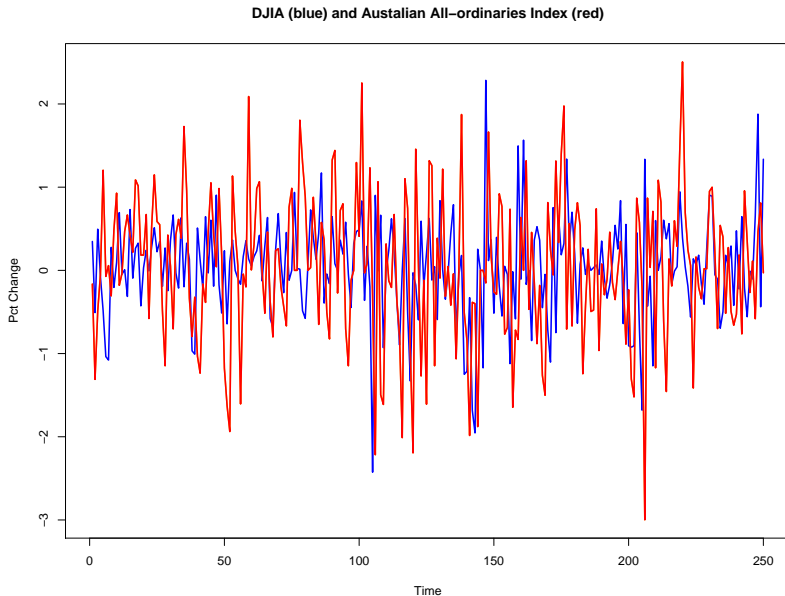
$$\Gamma(t+h, t) = \text{Cov}(\mathbf{X}_{t+h}, \mathbf{X}_t) \begin{bmatrix} \text{Cov}(X_{t+h,1} X_{t1}) & \text{Cov}(X_{t+h,1} X_{t2}) \\ \text{Cov}(X_{t+h,2} X_{t1}) & \text{Cov}(X_{t+h,2} X_{t2}) \end{bmatrix}$$

# Time Series Analysis

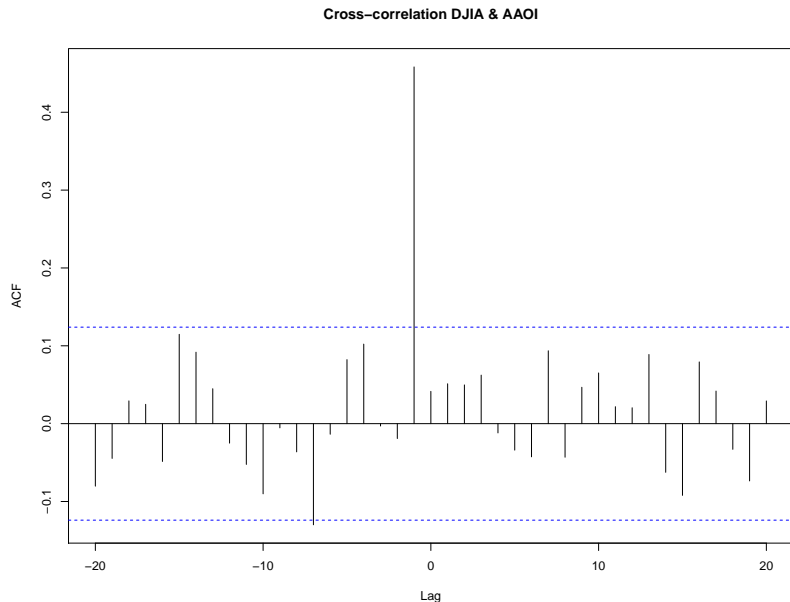
## Fitting a bivariate model DJIA and AAOI

```
ts.plot(djaopc2$DJIA, type="l",  
        main = "DJIA (blue) and Austalian  
        All-ordinaries Index (red)",  
        col="blue", lwd=2,  
        ylab="Pct Change",  
        ylim = c(range(djaopc2$AAOI)))  
lines(djaopc2$AAOI, col="red", lwd=2)  
tmp <- ccf(djaopc2$DJIA, djaopc2$AAOI,  
           main="Cross-correlation DJIA & AAOI")
```

# Time Series Analysis



# Time Series Analysis





# Time Series Analysis

```
dj.ar <- ar(djaopc2, order.max = 10)
dj.ar
```

```
Call:
ar(x = djaopc2, order.max = 10)
```

```
$ar
, , 1
```

	DJIA	AAOI
DJIA	-0.0148	0.0357
AAOI	0.6589	0.0998

dj.ar\$aic	0	1	2	3	4
	54.94030	0.00000	1.68532	8.45426	13.10385

# Time Series Analysis

In the R output the matrix labeled **\$ar** gives us the AR coefficients ( $\phi_{ij}$ ) for the bivariate model

$$\begin{bmatrix} Y_{t1} \\ Y_{t2} \end{bmatrix} = \begin{bmatrix} \phi_{11} & \phi_{12} \\ \phi_{21} & \phi_{22} \end{bmatrix} \begin{bmatrix} Y_{t-1,1} \\ Y_{t-1,2} \end{bmatrix} + \begin{bmatrix} Z_{t1} \\ Z_{t2} \end{bmatrix}$$

# Time Series Analysis

## Exponentially Weighted Moving Average

- ▶ Moving averages are often used to smooth volatility in data display an overall trend
- ▶ For example we may see a 60-day moving average overlaid on a chart of daily closing stock prices
- ▶ These are typically equally weighted so that the value for each day contributes equally
- ▶ Equally weighted moving averages are sensitive to extreme values

# Time Series Analysis

- ▶ To mitigate the effects of extremes we can weight the observations differently often putting more weight on the most recent observations
- ▶ The **exponentially weighted moving average**, EWMA, does this by using weights that decay exponentially
- ▶ An  $n$ -period EWMA of a series  $\{X_t\}$  is

$$\tilde{\mu}_t(n) = \sum_{i=0}^{n-1} w_i y_{t-i}, \quad w_i = \frac{\lambda^{i-1}}{\sum_{i=0}^{n-1} \lambda^{i-1}}$$

where  $0 < \lambda < 1$  is the decay parameter and  $y_t$  is the observation at time  $t$

- ▶ The closer  $\lambda$  is to 1, the more weight on the more recent observations

# Time Series Analysis

- ▶ As  $n \rightarrow \infty$ ,  $\lambda^n \rightarrow 0$ ,  $w_n \rightarrow 0$  and  $\tilde{\mu}_t(n)$  approaches

$$\tilde{\mu}_t(\lambda) = (1 - \lambda) \sum_{i=0}^{\infty} \lambda^i y_{t-i}$$

- ▶ EWMA can be computed recursively

$$\tilde{\mu}_t(\lambda) = (1 - \lambda)y_t + \lambda\tilde{\mu}_{t-1}(\lambda)$$

- ▶ One often uses the first value of the series or an average over some small window as a starting value for  $\tilde{\mu}_0(\lambda)$

# Time Series Analysis

## GARCH

- ▶ ARMA and ARIMA models focus on modeling and predicting the conditional mean of a time series
- ▶ In some areas, especially financial markets, it is important to model the volatility of the series
- ▶ Volatility is key to risk management, portfolio selection, and pricing derivatives, for example
- ▶ The class of models for time varying volatility is the generalized autoregressive heteroscedasticity (GARCH) models (also spelled *heteroskedasticity*)

# Time Series Analysis

- ▶ We can model the serial correlation in squared returns using a simple  $AR(p)$  model
- ▶ Assuming the returns are stationary and have *iid* errors with mean zero
- ▶ And the assumed variance is time dependent

$$\sigma_t^2 = a_0 + a_1\epsilon_{t-1}^2 + \cdots + a_p\epsilon_{t-p}^2$$

- ▶ We have an  $ARCH(p)$  model and we can test for ARCH effects

# Time Series Analysis

- ▶ If the ARCH test is significant we can model the volatility but this is impractical
- ▶ The GARCH( $p, q$ ) model extends ARCH( $p$ ), and when  $q = 0$  we have the latter
- ▶ Under GARCH( $p, q$ ) the conditional variance  $\sigma_t^2$  depends on the squared residuals in the  $p$  previous periods, and the conditional variance of the  $q$  previous periods

$$y_t = c + \epsilon_t, \quad \sigma_t^2 = a_0 + \sum_{i=1}^p a_i \epsilon_{t-i}^2 + \sum_{j=1}^q b_j \sigma_{t-j}^2$$

- ▶ For financial time series data, a GARCH(1, 1) model is usually adequate




# Time Series Analysis


```
par(mfrow=c(1,2))  
plot(spx.ts, type="l", col = 'blue', lwd =2,  
      ylab = "SPX Close",  
      main="S&P 500 Daily Close")  
  
spx.ret <- getReturns(spx.ts)  
plot(spx.ret, type="l", col = 'red', lwd =2,  
      ylab = "SPX logReturns",  
      main="Daily Returns")  
par(mfrow=c(1,1))
```

# Time Series Analysis

S&P 500 Daily Close



Daily Returns




# Time Series Analysis


```
par(mfrow=c(1,2))
tmp <- acf(spx.ret,
           main="ACF of SPX Daily Returns")
tmp <- acf(spx.ret^2,
           main="ACF of Squared SPX Daily Returns")
par(mfrow=c(1,1))
```

# Time Series Analysis

ACF of SPX Daily Returns



ACF of Squared SPX Daily Returns



# Time Series Analysis

```
library(MTS)
archTest(spx.ret)
```

Q(m) of squared series(LM test):

Test statistic: 7117 p-value: 0

Rank-based Test:

Test statistic: 7254.98 p-value: 0

**archTest** is significant so we move onto GARCH

# Time Series Analysis

```
library(tseries)
spx.gar <- garch(spx.ret)
summary(spx.gar)
```

# Time Series Analysis

Model:

GARCH(1,1)

Residuals:

Min	1Q	Median	3Q	Max
-10.9355	-0.5108	0.0565	0.6222	6.5912

Coefficient(s):

	Estimate	Std. Error	t value	Pr(> t )	
a0	8.50e-07	5.37e-08	15.8	<2e-16	***
a1	8.80e-02	1.53e-03	57.4	<2e-16	***
b1	9.09e-01	1.79e-03	508.4	<2e-16	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1

# Time Series Analysis

Based on the above results our model for the conditional variance is

$$\sigma_t^2 = 0.00000085 + 0.088\epsilon_{t-1}^2 + 0.909\sigma_{t-1}^2$$

The variance at time  $t$  relies on the squared errors from the prior period plus the variance from the prior period



# References

- Brockwell, P. J. and Davis, R. A. (1991).  
*Time Series: Theory and Methods*.  
Springer-Verlag, New York, second edition.
- Brockwell, P. J. and Davis, R. A. (1996).  
*Introduction to Time Series and Forecasting*.  
Springer-Verlag, New York.
- Carmona, R. A. (2004).  
*Statistical Analysis of Financial Data in S-PLUS*.  
Springer-Verlag.
- Wu, J. and Coggeshall, S. (2012).  
*Foundations of Predictive Analytics*.  
CRC Press, Boca Raton, FL.
- Zivot, E. and Wang, J. (2006).  
*Modeling Financial Time Series with S-PLUS*.  
Springer-Verlag, New York, second edition.